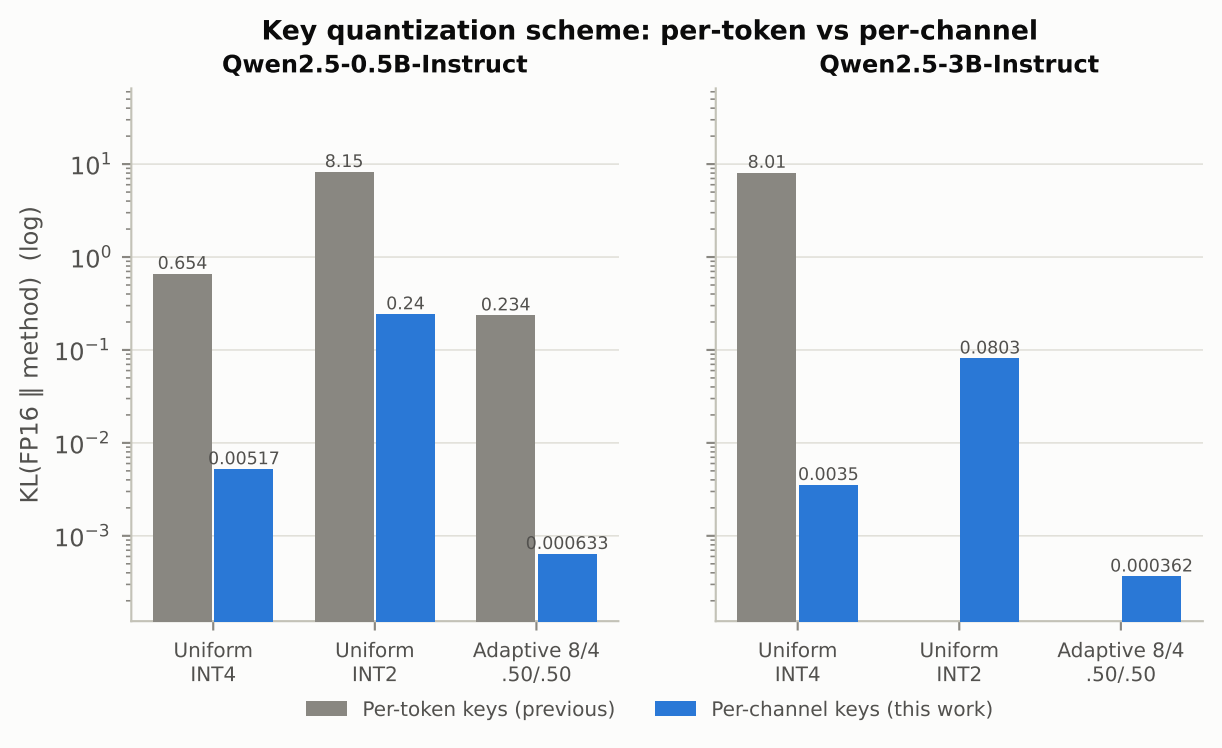
Key quantization scheme: per-token vs per-channel
Qwen2.5-0.5B-Instruct
Qwen2.5-3B-Instruct
KL(FP16 ‖ method) (log)
10^{1}
10^{0}
10^{-1}
10^{-2}
10^{-3}
0.654
0.00517
8.15
0.24
0.234
0.000633
8.01
0.0035
0.0803
0.000362
Uniform INT4
Uniform INT2
Adaptive 8/4 .50/.50
Per-token keys (previous)
Per-channel keys (this work)